

Jordan Kail

Staff Software Engineer — AI, Agents, Data & Machine Learning

Denver, CO | jckail13@gmail.com | linkedin.com/in/jckail | github.com/jckail | jordan-kail.com

Summary

With over 13 years of experience in AI, analytics and machine learning, I now lead the agents platform and data engineering at Together AI, building the agent harness, internal agents factory and proprietary IP behind them. Before that I built AI classifiers at Meta and model-governance frameworks at Prove, and turned data into impact for startups and Fortune 50 companies alike.

Experience

Together AI — Staff Software Engineer — San Francisco, CA

02/2025 – Present

- Tech lead for data engineering since joining: grew the team from just myself to 15+ engineers while setting technical direction for the data platform behind Together's AI acceleration cloud.
- Tech lead for the agents platform: created the agent harness, built an internal agents factory and used it to deliver production agents, and contributed to multiple patent-pending innovations in agent systems.
- Built and operate agents for infrastructure automation, finance automation and other internal workflows that augment the teams they serve, extending team capacity while improving on-call response and infrastructure maintenance.
- Build the data platform behind Together's AI acceleration cloud: the pipelines, storage, and telemetry that turn inference and training traffic into product, reliability, and capacity signals.
- Build tracing, replay and evaluation harnesses so internal teams can measure and regression-test LLM agents before they ship.

Prove Identity — Staff Software Engineer - Data — Denver, CO

06/2023 – 01/2025

- Built agents and harnesses that govern the statistical models used for identity and fraud resolution; these became core frameworks.
- Built AI-driven Retrieval-Augmented Generation (RAG) chatbots with Airflow, LangChain, and OpenAI, automating 150+ human-hours weekly.
- Led company-wide migration from legacy Java/Oracle infrastructure to cloud-native Go/PostgreSQL stack, improving API response times from 40s to 12ms and slashing operational costs by 95%.
- Deployed an event-driven data streaming platform with Go, Kafka, and Flink, transitioning 1,200+ batch jobs to real-time processing. Reduced data availability lag from days to seconds.

Sabbatical — Career break (digital nomad) — Location independent

10/2022 – 05/2023

- Stepped away from corporate life to live a digital nomad life with my wife.

Meta (Facebook) — Senior Data Engineer — Seattle, WA; Menlo Park, CA; Remote, USA

01/2021 – 09/2022

- Built AI/ML classifiers and the data, feature and evaluation pipelines around them.
- Designed and deployed 100+ ML pipelines using Airflow, Spark, and PyTorch, integrating NLP and computer vision models. Improved ad targeting precision and personalized notifications, driving higher engagement across billions of users.
- Lead data engineer for Facebook Public Groups, Community Chats, and cross-platform initiatives, leading 10+ engineers.
- Engineered an automated framework for generating thousands of asynchronous Spark data pipelines, increasing compute efficiency by 66%.
- Re-architected petabyte-scale data models, reducing storage costs by 25%.

Deloitte — Consultant - AI & Advanced Analytics — Atlanta, GA; Menlo Park, CA

12/2018 – 12/2020

- Automated approximately 31% of human processed healthcare claims using transformer machine learning models, saving 250,000+ hours annually.
- Optimized exabyte-scale video reliability metrics, reducing daily processing time by 90% while expanding metric coverage.
- Led teams of 20+ consultants on high-profile Fortune 50 engagements, delivering advanced technical solutions aligned with client needs.

Wide Open West — Senior Data Engineer — Denver, CO

11/2017 – 12/2018

- Built Machine Learning applications using custom classification and churn models, driving a 22% YoY increase in customer package upgrades.
- Led a team of 5 data practitioners, providing BI and data insights to sales, product, and engineering teams company-wide.

Common Spirit Health — Data Engineer — Denver, CO

09/2016 – 11/2017

- Architected and delivered new rest APIs and data lakes, improving data processing time for external partner data products from 7 days to 5 minutes.
- Developed machine learning pipelines using Python and SQL to forecast hospital procedures, billing, and staffing. Reduced billing turnaround from 14 to 7 days.

AcuStream (R1) — Software Engineer - Data — Boulder, CO

04/2013 – 09/2016

- Built a custom invoicing system leveraging rule-based algorithms and machine learning, driving over \$300M in annual revenue.
- Migrated clients from SFTP to real-time APIs with Flask and Apache Kafka, reducing data delivery times by 50%.

Skills

AI & Machine Learning: Amazon Bedrock, Chroma, Claude AI, Google Gemini, Hugging Face, Keras, LangChain, Mistral AI, OpenAI, PyTorch, TensorFlow, Vertex AI, Agent Platforms, Agent Harnesses, Agent Evaluation, LLM Evaluation, AI Classifiers, Model Governance, RAG

Programming Languages: C++, Go, Java, JavaScript, PHP, Python, Rust, Scala, SQL, TypeScript

Data Engineering: Airflow, Kafka, Pulsar, Cloud Composer, Cloud Data Fusion, Dataprep, Dataproc, DBT, Flink, Kinesis Firehose, PubSub

Big Data: Hadoop, Iceberg, Spark, BigQuery, Databricks, Presto, Snowflake, Trino

Data Science: Jupyter, Looker, NumPy, Pandas, Scikit-learn, Streamlit, Tableau

Databases: Cassandra, DuckDB, DynamoDB, Elasticsearch, Firestore, MongoDB, MySQL, Neo4j, Pinecone, PostgreSQL, Redis, RocksDB, Supabase

Cloud Computing: Azure, AWS, Google Cloud

DevOps: Docker, Git, GitHub Actions, Grafana, Jaeger, Kubernetes, NGINX, OpenTelemetry, Prometheus, Splunk, Terraform

Web Development: CSS3, Django, FastAPI, Flask, Gin, GraphQL, HTML5, Material UI, Node.js, React, Redux, SQLAlchemy, Tailwind CSS, Vite, WebAssembly

Development Tools: Figma, Linux, Playwright, Selenium, Retool

Projects

Super Teacher — github.com/jckail/superteacher, jckail.com/superteacher

A classroom copilot that answers 'who needs me today, and why?' with a ranked needs-attention list, a gradebook, attendance, and a streaming Claude assistant grounded in the class data.

Jobbr — github.com/jckail/Jobbr, jckail.com/jobbr

An AI-assisted job-search prototype that parses tech-company postings into structured data and matches them to resumes, using Python, LangChain, FastAPI and OpenAI.

goPilot — github.com/jckail/goPilot, jckail.com/gopilot

Agentic AI developer assistant and cli tool to help debug, develop, and compile apps written in GO.

AI Billing System — github.com/jckail/ai_chat_billing_app, jckail.com/aibilling

A proof-of-concept application for tracking and billing AI chat thread interactions with real-time cost metrics and analytics.

Portfolio Website — github.com/jckail/portfolio

A fully custom-built portfolio site showcasing my skills and work.

PointUp — github.com/jckail/point_bot, pointup.io

Track airline miles, hotel points, and credit card rewards in one place.

Join Group via QR — techcrunch.com

Featured on TechCrunch - Won internal hackathon and created the ability for Facebook group admins to invite users to their groups by generating a QR Code.